

# Project 6-b

September 16, 2025

## 1 Customer Churn Prediction

### 1.1 Task: Data Exploration and Preprocessing Report

#### 1.1.1 Task 1: Import Necessary Libraries and Load the Data

```
[7]: # Importing standard libraries for data manipulation and visualization
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Setting visualization style
sns.set_style("whitegrid")
%matplotlib inline

# Load the dataset
# Note: Please ensure the file path is correct. We'll assume the file is named ↴
# 'customer_data.csv'
df = pd.read_csv('Customer_data - customer_data.csv')
```

#### 1.1.2 Task 2: Initial Data Exploration

```
[8]: # 2.1 Display the first 5 rows to peek at the data
print("First 5 rows of the dataset:")
display(df.head())

# 2.2 Get the dimensions of the dataset (rows, columns)
print(f"\nDataset shape: {df.shape}")

# 2.3 Get basic info about data types and non-null counts
print("\nDataset info:")
df.info()
```

First 5 rows of the dataset:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
0	7590-VHVEG	Female		0	Yes	No	1	No
1	5575-GNVDE	Male		0	No	No	34	Yes

```

2 3668-QPYBK   Male          0    No     No    2    Yes
3 7795-CFOCW   Male          0    No     No    45   No
4 9237-HQITU   Female        0    No     No    2    Yes

```

```

      MultipleLines InternetService OnlineSecurity ... DeviceProtection \
0  No phone service           DSL       No ...      No
1            No                 DSL       Yes ...     Yes
2            No                 DSL       Yes ...      No
3  No phone service           DSL       Yes ...     Yes
4            No      Fiber optic     No ...      No

```

```

TechSupport StreamingTV StreamingMovies      Contract PaperlessBilling \
0        No        No        No Month-to-month      Yes
1        No        No        No One year          No
2        No        No        No Month-to-month      Yes
3      Yes        No        No One year          No
4        No        No        No Month-to-month      Yes

```

```

      PaymentMethod MonthlyCharges  TotalCharges  Churn
0  Electronic check        29.85      29.85      No
1  Mailed check           56.95    1889.50      No
2  Mailed check           53.85     108.15     Yes
3  Bank transfer (automatic)  42.30    1840.75      No
4  Electronic check        70.70     151.65     Yes

```

[5 rows x 21 columns]

Dataset shape: (7043, 21)

Dataset info:

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #  Column          Non-Null Count Dtype  
--- 
 0  customerID      7043 non-null  object  
 1  gender          7043 non-null  object  
 2  SeniorCitizen   7043 non-null  int64   
 3  Partner          7043 non-null  object  
 4  Dependents      7043 non-null  object  
 5  tenure           7043 non-null  int64   
 6  PhoneService     7043 non-null  object  
 7  MultipleLines    7043 non-null  object  
 8  InternetService  7043 non-null  object  
 9  OnlineSecurity   7043 non-null  object  
 10  OnlineBackup     7043 non-null  object  
 11  DeviceProtection 7043 non-null  object 

```

```

12 TechSupport      7043 non-null  object
13 StreamingTV      7043 non-null  object
14 StreamingMovies   7043 non-null  object
15 Contract         7043 non-null  object
16 PaperlessBilling 7043 non-null  object
17 PaymentMethod    7043 non-null  object
18 MonthlyCharges   7043 non-null  float64
19 TotalCharges     7032 non-null  float64
20 Churn            7043 non-null  object
dtypes: float64(2), int64(2), object(17)
memory usage: 1.1+ MB

```

### 1.1.3 Task 3: Check for Missing Values

```
[11]: # 3.1 Check for missing values in each column
print("Missing values per column:")
missing_values = df.isnull().sum()
print(missing_values[missing_values > 0]) # Only show columns with missing values

# If no output from the above line, confirm no missing values exist
if missing_values.sum() == 0:
    print("No missing values detected in any column.")
else:
    print(f"\nTotal missing values in dataset: {missing_values.sum()}")
```

Missing values per column:  
TotalCharges 11  
dtype: int64

Total missing values in dataset: 11

### 1.1.4 Task 4: Deep Dive into Categorical Variables

```
[13]: # 4.1 List all categorical columns
categorical_columns = df.select_dtypes(include='object').columns.tolist()
print("Categorical columns:", categorical_columns)

# 4.2 Analyze unique values in each categorical column
print("\nUnique values in categorical columns:")
for col in categorical_columns:
    print(f"{col}: {df[col].nunique()} unique values -> {df[col].unique()}")
    #print(df[col].value_counts()) # Uncomment this line for a detailed count of each category

# 4.3 Visualize the distribution of the target variable 'Churn'
plt.figure(figsize=(6, 4))
sns.countplot(x=df['Churn'])
```

```

plt.title('Distribution of Target Variable: Churn')
plt.show()

# Calculate the churn rate
churn_rate = (df['Churn'].value_counts(normalize=True) * 100)[['Yes']]
print(f"\nChurn Rate: {churn_rate:.2f}%")

```

Categorical columns: ['customerID', 'gender', 'Partner', 'Dependents',  
'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity',  
'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',  
'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod', 'Churn']

Unique values in categorical columns:

```

customerID: 7043 unique values -> ['7590-VHVEG' '5575-GNVDE' '3668-QPYBK' ...
'4801-JZAZL' '8361-LTMKD'
'3186-AJIEK']

gender: 2 unique values -> ['Female' 'Male']

Partner: 2 unique values -> ['Yes' 'No']

Dependents: 2 unique values -> ['No' 'Yes']

PhoneService: 2 unique values -> ['No' 'Yes']

MultipleLines: 3 unique values -> ['No phone service' 'No' 'Yes']

InternetService: 3 unique values -> ['DSL' 'Fiber optic' 'No']

OnlineSecurity: 3 unique values -> ['No' 'Yes' 'No internet service']

OnlineBackup: 3 unique values -> ['Yes' 'No' 'No internet service']

DeviceProtection: 3 unique values -> ['No' 'Yes' 'No internet service']

TechSupport: 3 unique values -> ['No' 'Yes' 'No internet service']

StreamingTV: 3 unique values -> ['No' 'Yes' 'No internet service']

StreamingMovies: 3 unique values -> ['No' 'Yes' 'No internet service']

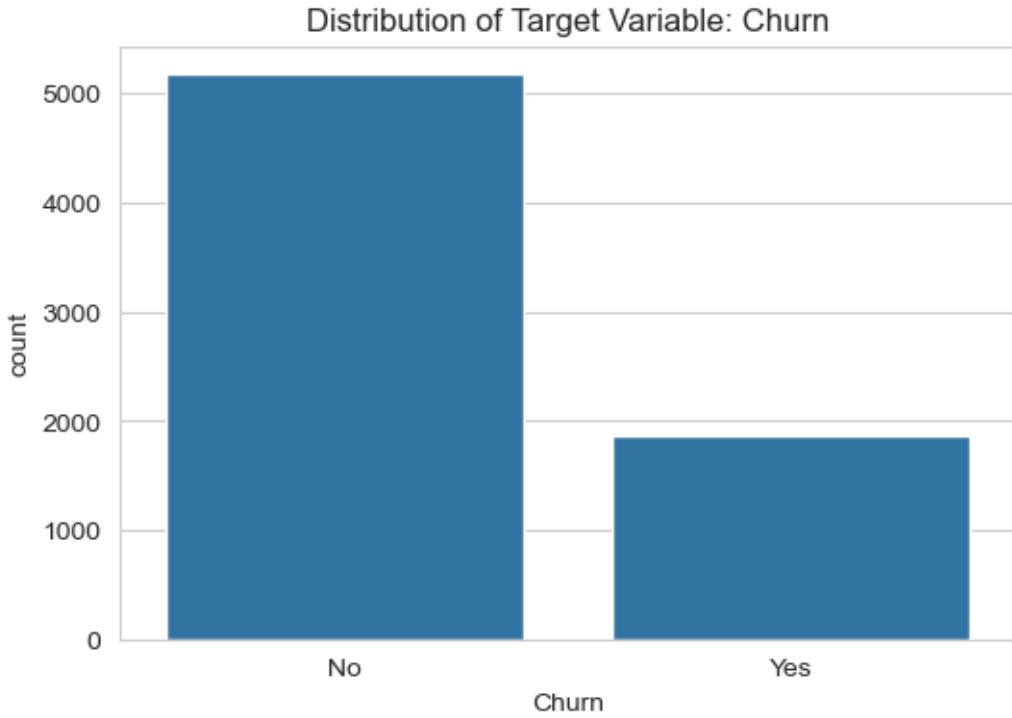
Contract: 3 unique values -> ['Month-to-month' 'One year' 'Two year']

PaperlessBilling: 2 unique values -> ['Yes' 'No']

PaymentMethod: 4 unique values -> ['Electronic check' 'Mailed check' 'Bank
transfer (automatic)'
'Credit card (automatic)']

Churn: 2 unique values -> ['No' 'Yes']

```



Churn Rate: 26.54%

### 1.1.5 Task 5: Deep Dive into Numerical Variables

```
[15]: # 5.1 List all numerical columns
numerical_columns = ['tenure', 'MonthlyCharges', 'TotalCharges', ↴
                     'SeniorCitizen']
print("Numerical columns:", numerical_columns)

# 5.2 Generate descriptive statistics
print("\nDescriptive statistics for numerical columns:")
display(df[numerical_columns].describe())

# 5.3 Plot distributions for numerical columns
fig, axes = plt.subplots(2, 2, figsize=(14, 10))
fig.suptitle('Distribution of Numerical Variables', fontsize=16)
axes = axes.ravel() # Flatten the 2x2 array of axes

for i, col in enumerate(numerical_columns):
    df[col].hist(ax=axes[i], bins=30)
    axes[i].set_title(col)
    axes[i].set_ylabel('Count')
```

```

plt.tight_layout()
plt.show()

# 5.4 Check for outliers using boxplots
fig, axes = plt.subplots(2, 2, figsize=(14, 10))
fig.suptitle('Boxplots of Numerical Variables (Outlier Check)', fontsize=16)
axes = axes.ravel()

for i, col in enumerate(numerical_columns):
    sns.boxplot(y=df[col], ax=axes[i])
    axes[i].set_title(col)

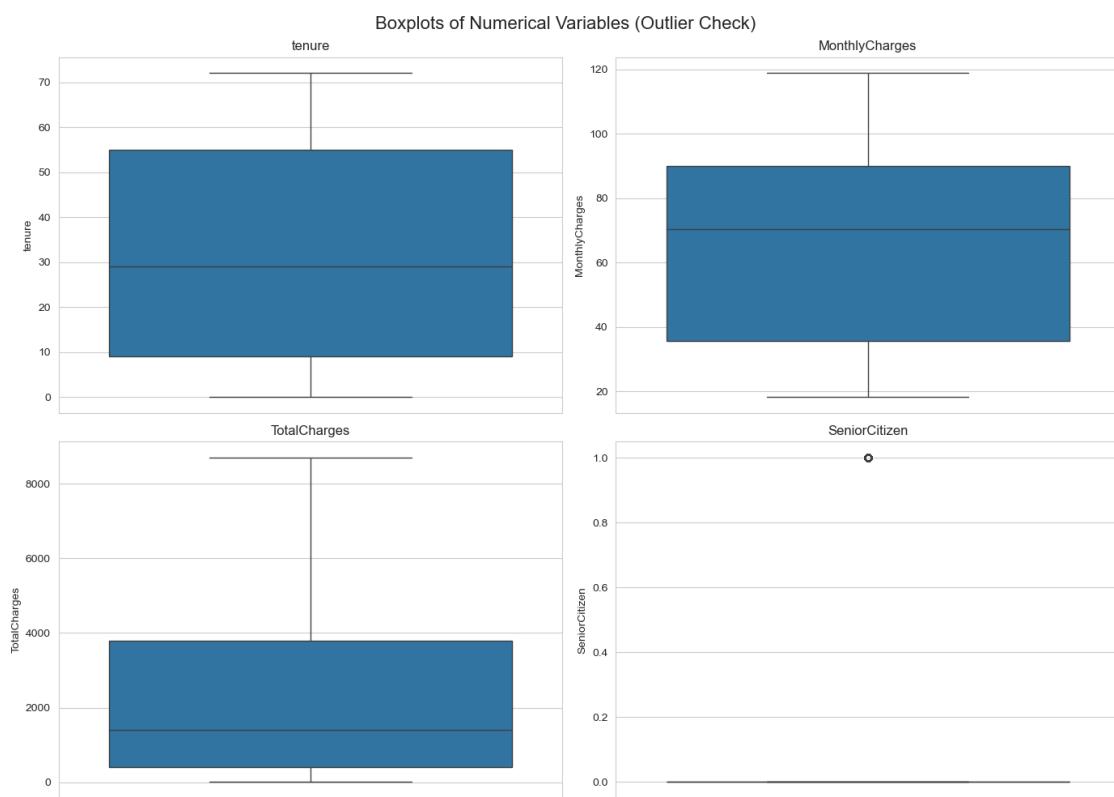
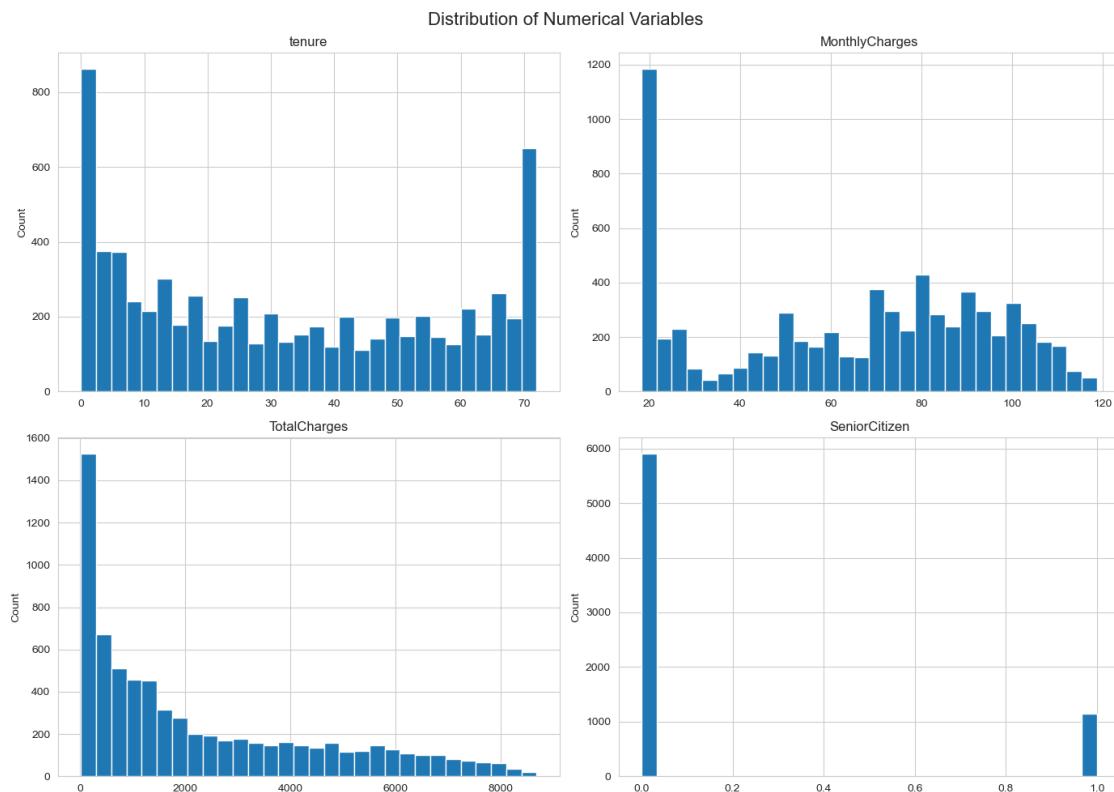
plt.tight_layout()
plt.show()

```

Numerical columns: ['tenure', 'MonthlyCharges', 'TotalCharges', 'SeniorCitizen']

Descriptive statistics for numerical columns:

	tenure	MonthlyCharges	TotalCharges	SeniorCitizen
count	7043.000000	7043.000000	7032.000000	7043.000000
mean	32.371149	64.761692	2283.300441	0.162147
std	24.559481	30.090047	2266.771362	0.368612
min	0.000000	18.250000	18.800000	0.000000
25%	9.000000	35.500000	401.450000	0.000000
50%	29.000000	70.350000	1397.475000	0.000000
75%	55.000000	89.850000	3794.737500	0.000000
max	72.000000	118.750000	8684.800000	1.000000



### 1.1.6 Task 6: Bivariate Analysis - Relationship with Churn

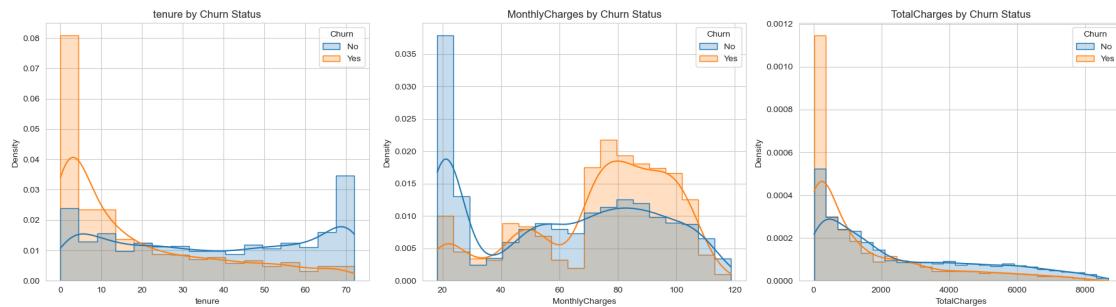
```
[18]: # 6.1 Analyze numerical variables vs. Churn
fig, axes = plt.subplots(1, 3, figsize=(18, 5))
num_vars = ['tenure', 'MonthlyCharges', 'TotalCharges']

for i, var in enumerate(num_vars):
    sns.histplot(data=df, x=var, hue='Churn', kde=True, element='step', u
    ↪stat='density', common_norm=False, ax=axes[i])
    axes[i].set_title(f'{var} by Churn Status')

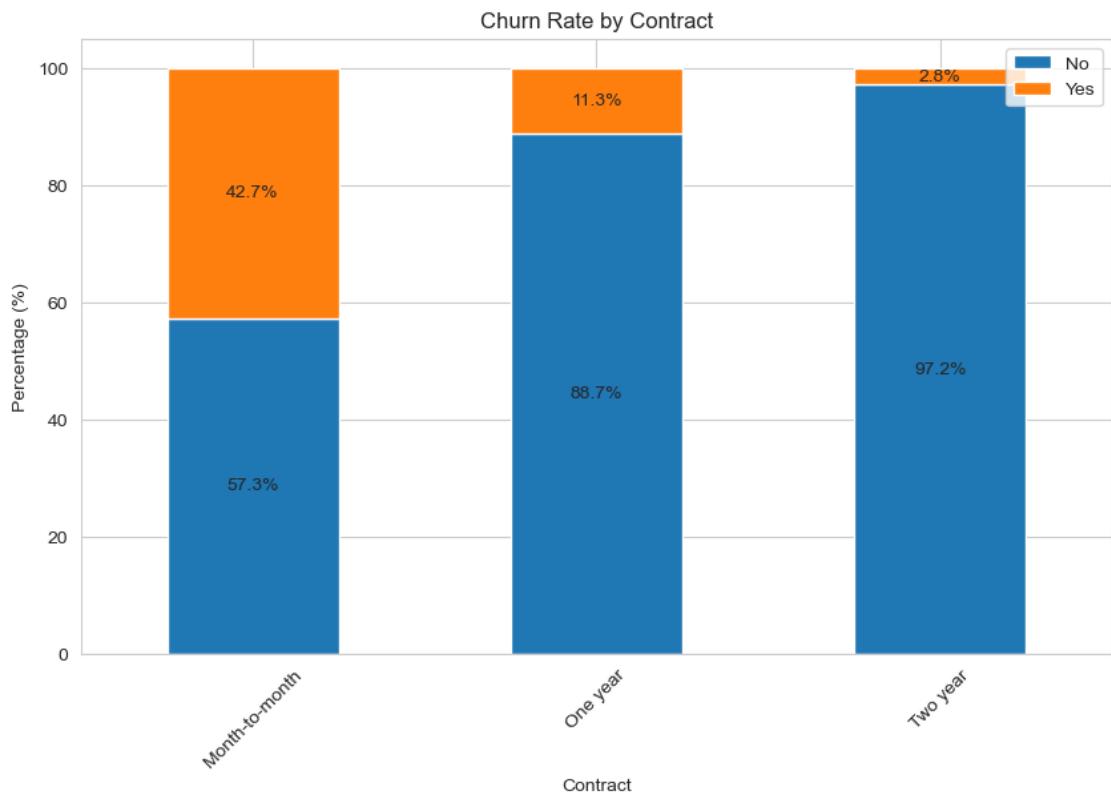
plt.tight_layout()
plt.show()

# 6.2 Analyze key categorical variables vs. Churn
key_cat_vars = ['Contract', 'InternetService', 'PaymentMethod', 'TechSupport']

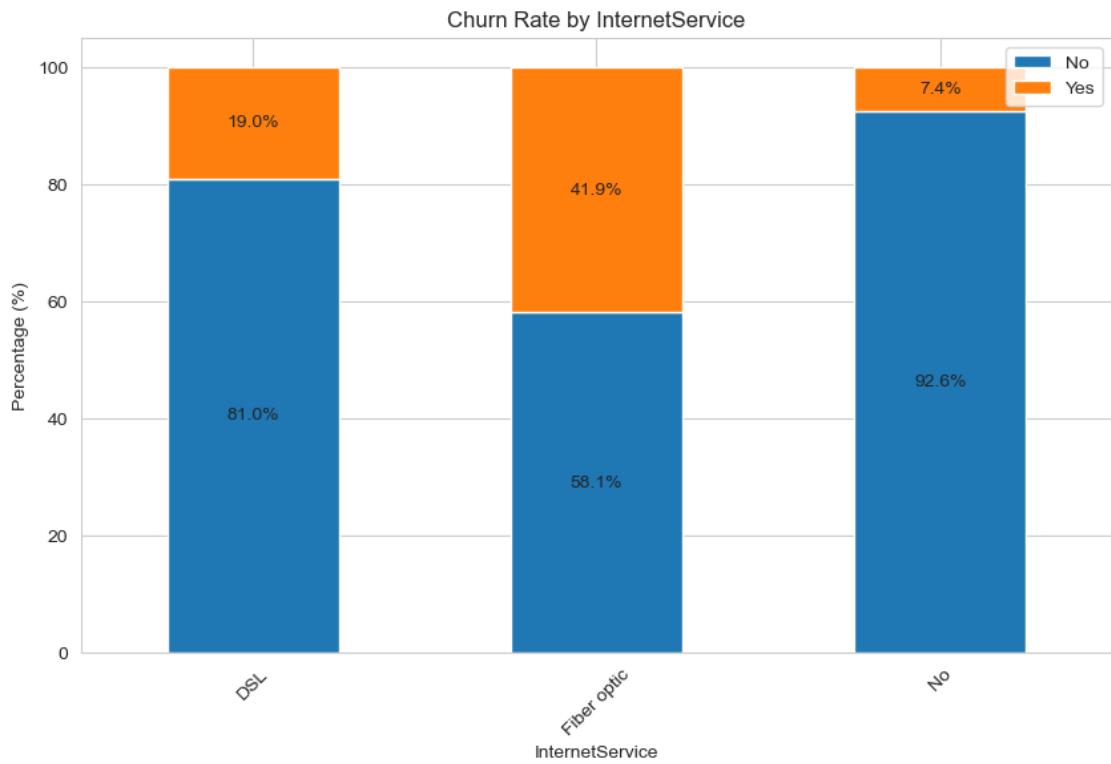
for col in key_cat_vars:
    plt.figure(figsize=(10, 5))
    # Create a proportional crosstab
    prop_churn = pd.crosstab(df[col], df['Churn'], normalize='index') * 100
    ax = prop_churn.plot(kind='bar', stacked=True, figsize=(10, 6))
    plt.title(f'Churn Rate by {col}')
    plt.ylabel('Percentage (%)')
    plt.xticks(rotation=45)
    # Add percentage labels on the bars
    for c in ax.containers:
        labels = [f'{v.get_height():.1f}%' if v.get_height() > 0 else '' for v
        ↪in c]
        ax.bar_label(c, labels=labels, label_type='center')
    plt.legend(loc='upper right')
    plt.show()
```



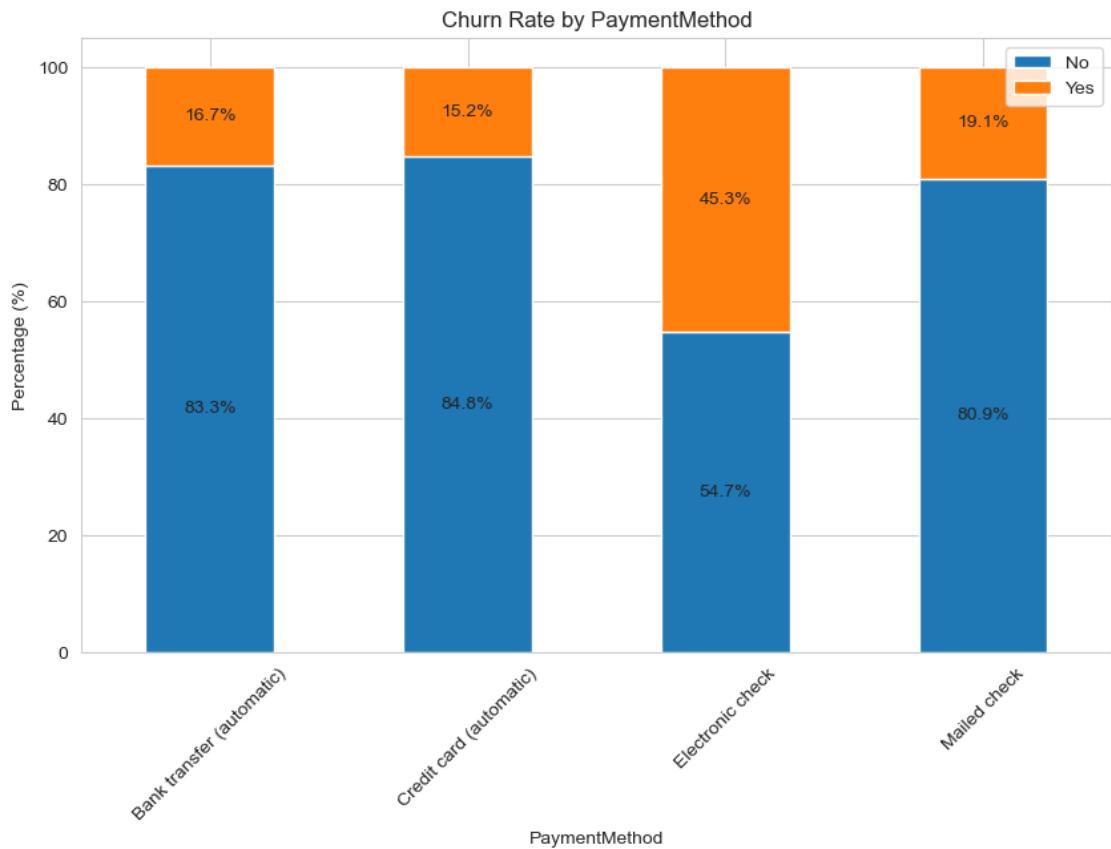
<Figure size 1000x500 with 0 Axes>



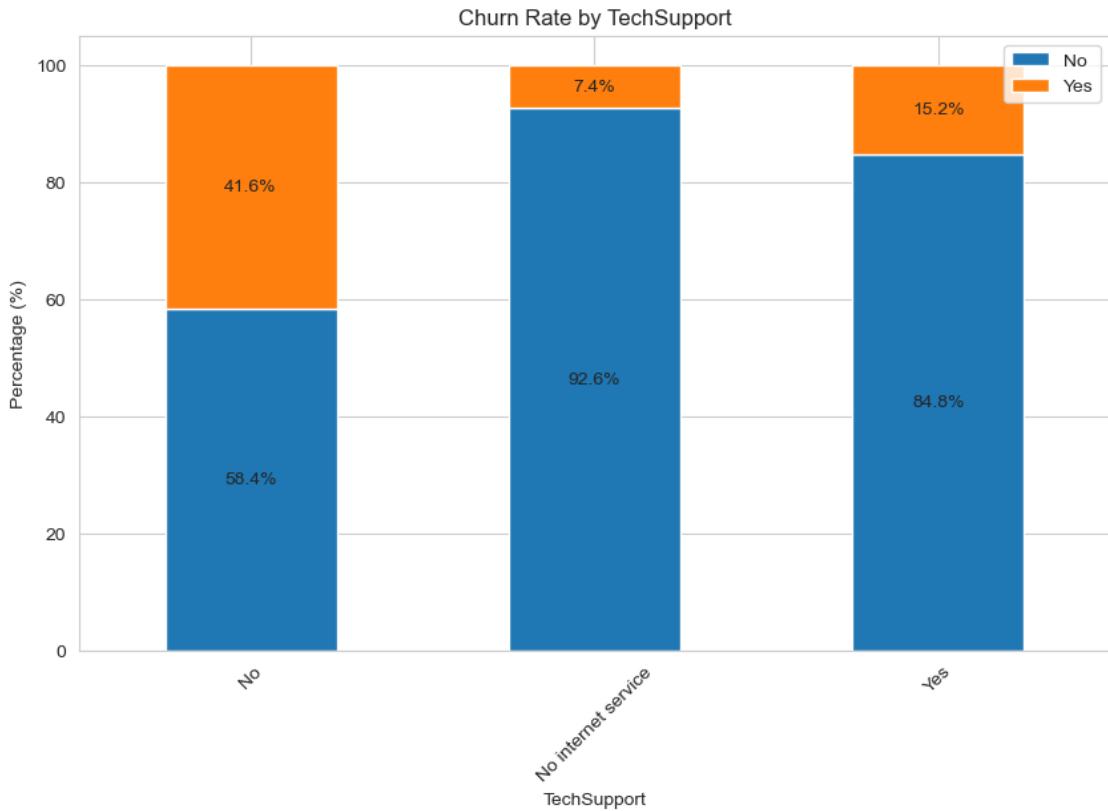
<Figure size 1000x500 with 0 Axes>



<Figure size 1000x500 with 0 Axes>



<Figure size 1000x500 with 0 Axes>



### 1.1.7 Task 7: Data Preprocessing

```
[17]: # 7.1 Handle missing values (if any were found in Task 3)
# Example code if 'TotalCharges' had missing values (though our check showed ↴none):
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce') # Force ↴conversion, making invalid parsing NaN
df['TotalCharges'].fillna(0, inplace=True) # Fill NaN with 0 (logical for new ↴customers)

# 7.2 Encode the target variable 'Churn'
df['Churn'] = df['Churn'].apply(lambda x: 1 if x == 'Yes' else 0)

# 7.3 Separate features (X) and target (y)
X = df.drop('Churn', axis=1)
y = df['Churn']

# 7.4 Split into training and test set NOW, before any encoding, to avoid data ↴leakage
from sklearn.model_selection import train_test_split
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42, stratify=y)

print(f"Training set size: {X_train.shape}")
print(f"Test set size: {X_test.shape}")

# 7.5 Identify which columns need encoding
# Numerical columns: We will standardize them later
# Categorical columns: We will one-hot encode them
categorical_cols = X_train.select_dtypes(include='object').columns.tolist()
numerical_cols = X_train.select_dtypes(include=['int64', 'float64']).columns.
    ↪tolist()

print("Categorical columns to encode:", categorical_cols)
print("Numerical columns to scale:", numerical_cols)

# 7.6 Create a preprocessing pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.pipeline import Pipeline

# Preprocessor for numerical columns: Standardization
numerical_transformer = Pipeline(steps=[
    ('scaler', StandardScaler())
])

# Preprocessor for categorical columns: One-Hot Encoding
categorical_transformer = Pipeline(steps=[
    ('onehot', OneHotEncoder(drop='first', handle_unknown='ignore')) # ↪
    ↪ 'drop='first'' to avoid multicollinearity
])

# Bundle preprocessing for numerical and categorical data
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numerical_transformer, numerical_cols),
        ('cat', categorical_transformer, categorical_cols)
    ]
)

# 7.7 Fit the preprocessor on the training data and transform both sets
X_train_processed = preprocessor.fit_transform(X_train)
X_test_processed = preprocessor.transform(X_test)

# Get feature names after one-hot encoding
# This is a bit more complex with ColumnTransformer but possible
onehot_columns = preprocessor.named_transformers_['cat']['onehot'].
    ↪get_feature_names_out(categorical_cols)

```

```

all_feature_names = numerical_cols + list(onehot_columns)

print(f"\nFinal processed training set shape: {X_train_processed.shape}")
print(f"Processed feature names: {len(all_feature_names)}")
# print(all_feature_names) # Uncomment to see the full list of engineered
    ↴features

```

Training set size: (5634, 20)  
Test set size: (1409, 20)  
Categorical columns to encode: ['customerID', 'gender', 'Partner', 'Dependents',  
'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity',  
'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',  
'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod']  
Numerical columns to scale: ['SeniorCitizen', 'tenure', 'MonthlyCharges',  
'TotalCharges']

Final processed training set shape: (5634, 5663)  
Processed feature names: 5663

## 1.2 Task: Build a Machine Learning Model for Customer Churn Prediction

### 1.2.1 Step 1: Import Additional Modeling Libraries

```
[19]: # Import model algorithms
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier

# Import model evaluation tools
from sklearn.metrics import classification_report, confusion_matrix,
    ↴ConfusionMatrixDisplay
from sklearn.metrics import accuracy_score, precision_score, recall_score,
    ↴f1_score
from sklearn.metrics import roc_auc_score, roc_curve, precision_recall_curve

# Import hyperparameter tuning
from sklearn.model_selection import GridSearchCV

import warnings
warnings.filterwarnings('ignore')
```

### 1.2.2 Step 2: Establish a Baseline with Multiple Algorithms

```
[ ]: # Initialize a list of models to test
models = {
    'Logistic Regression': LogisticRegression(random_state=42,
        ↴class_weight='balanced', max_iter=1000),
    'Random Forest': RandomForestClassifier(random_state=42,
        ↴class_weight='balanced', n_estimators=100),
```

```

'Gradient Boosting': GradientBoostingClassifier(random_state=42,
    ↪n_estimators=100) # GB doesn't have class_weight, we'll handle it via tuning
}

# Dictionary to store baseline results
baseline_results = {}

# Train and evaluate each model
for name, model in models.items():
    # Fit the model on the training data
    model.fit(X_train_processed, y_train)

    # Predict on the test data
    y_pred = model.predict(X_test_processed)
    y_pred_proba = model.predict_proba(X_test_processed)[:, 1] # Get
    ↪probabilities for the positive class

    # Calculate key metrics
    accuracy = accuracy_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    roc_auc = roc_auc_score(y_test, y_pred_proba)

    # Store results
    baseline_results[name] = {
        'Accuracy': accuracy,
        'Precision': precision,
        'Recall': recall,
        'F1-Score': f1,
        'ROC-AUC': roc_auc
    }

    # Print a quick report
    print(f"--- {name} ---")
    print(f"ROC-AUC: {roc_auc:.4f}, F1-Score: {f1:.4f}\n")

# Create a DataFrame for easy comparison of baseline results
df_baseline = pd.DataFrame(baseline_results).T
print("Baseline Model Comparison:")
display(df_baseline.sort_values(by='ROC-AUC', ascending=False))

```

### 1.2.3 Step 3: Hyperparameter Tuning with Grid Search

```
[23]: # Define the parameter grid for Gradient Boosting
param_grid = {
    'n_estimators': [100, 200],
    'learning_rate': [0.05, 0.1],
    'max_depth': [3, 5],
    'subsample': [0.8, 1.0] # Fraction of samples used for fitting
}

# Initialize the model
gb_model = GradientBoostingClassifier(random_state=42)

# Set up GridSearchCV
print("Starting Grid Search for Gradient Boosting...")
grid_search = GridSearchCV(
    estimator=gb_model,
    param_grid=param_grid,
    scoring='roc_auc', # We want to maximize ROC-AUC
    cv=5, # 5-fold cross-validation
    n_jobs=-1, # Use all available CPU cores
    verbose=1
)

# Perform the grid search (this will take some time)
grid_search.fit(X_train_processed, y_train)
print("Grid Search Complete!")

# Print the best parameters and score
print(f"\nBest parameters found: {grid_search.best_params_}")
print(f"Best cross-validation ROC-AUC score: {grid_search.best_score_.:.4f}")

# Get the best model from the grid search
best_model = grid_search.best_estimator_
```

Starting Grid Search for Gradient Boosting...  
Fitting 5 folds for each of 16 candidates, totalling 80 fits  
Grid Search Complete!

Best parameters found: {'learning\_rate': 0.1, 'max\_depth': 3, 'n\_estimators': 100, 'subsample': 0.8}  
Best cross-validation ROC-AUC score: 0.8482

#### 1.2.4 Step 4: Final Evaluation on the Test Set

```
[25]: # Make final predictions using the best model
final_predictions = best_model.predict(X_test_processed)
final_predictions_proba = best_model.predict_proba(X_test_processed)[:, 1]

# 4.1 Generate a comprehensive classification report
print("### Final Classification Report ###")
print(classification_report(y_test, final_predictions, target_names=['Not Churn', 'Churn']))

# 4.2 Plot a Confusion Matrix
fig, ax = plt.subplots(figsize=(8, 6))
cm = confusion_matrix(y_test, final_predictions)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['Not Churn', 'Churn'])
disp.plot(cmap='Blues', ax=ax)
plt.title("Confusion Matrix - Tuned Gradient Boosting Model")
plt.show()

# 4.3 Plot ROC Curve and Precision-Recall Curve
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 6))

# ROC Curve
fpr, tpr, _ = roc_curve(y_test, final_predictions_proba)
roc_auc = roc_auc_score(y_test, final_predictions_proba)
ax1.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (AUC = {roc_auc:.2f})')
ax1.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
ax1.set_xlim([0.0, 1.0])
ax1.set_ylim([0.0, 1.05])
ax1.set_xlabel('False Positive Rate')
ax1.set_ylabel('True Positive Rate')
ax1.set_title('Receiver Operating Characteristic (ROC) Curve')
ax1.legend(loc="lower right")

# Precision-Recall Curve
precision_curve, recall_curve, _ = precision_recall_curve(y_test, final_predictions_proba)
avg_precision = np.average(precision_curve)
ax2.plot(recall_curve, precision_curve, color='green', lw=2, label=f'Avg Precision = {avg_precision:.2f}')
ax2.set_xlim([0.0, 1.0])
ax2.set_ylim([0.0, 1.05])
ax2.set_xlabel('Recall')
ax2.set_ylabel('Precision')
ax2.set_title('Precision-Recall Curve')
```

```

ax2.legend(loc="lower left")

plt.tight_layout()
plt.show()

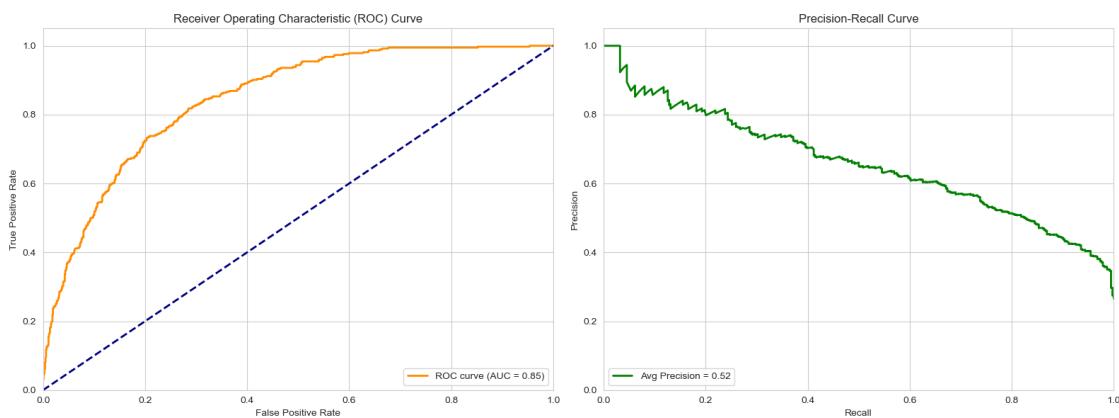
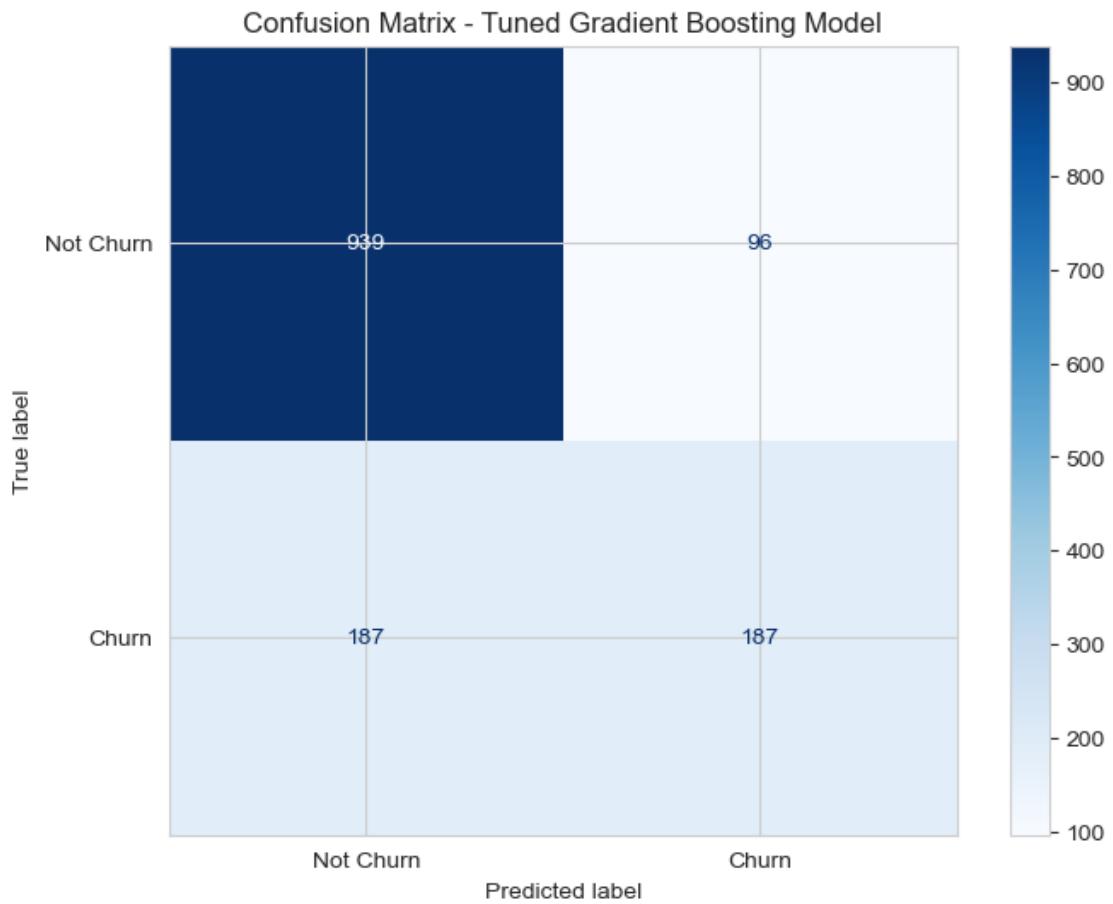
# Print final metrics
final_accuracy = accuracy_score(y_test, final_predictions)
final_recall = recall_score(y_test, final_predictions)
final_precision = precision_score(y_test, final_predictions)
final_f1 = f1_score(y_test, final_predictions)

print(f"\nFinal Model Performance on Test Set:")
print(f"Accuracy: {final_accuracy:.4f}")
print(f"Precision: {final_precision:.4f}")
print(f"Recall: {final_recall:.4f}")
print(f"F1-Score: {final_f1:.4f}")
print(f"ROC-AUC: {roc_auc:.4f}")

```

### ### Final Classification Report ###

	precision	recall	f1-score	support
Not Churn	0.83	0.91	0.87	1035
Churn	0.66	0.50	0.57	374
accuracy			0.80	1409
macro avg	0.75	0.70	0.72	1409
weighted avg	0.79	0.80	0.79	1409



Final Model Performance on Test Set:

Accuracy: 0.7991

Precision: 0.6608

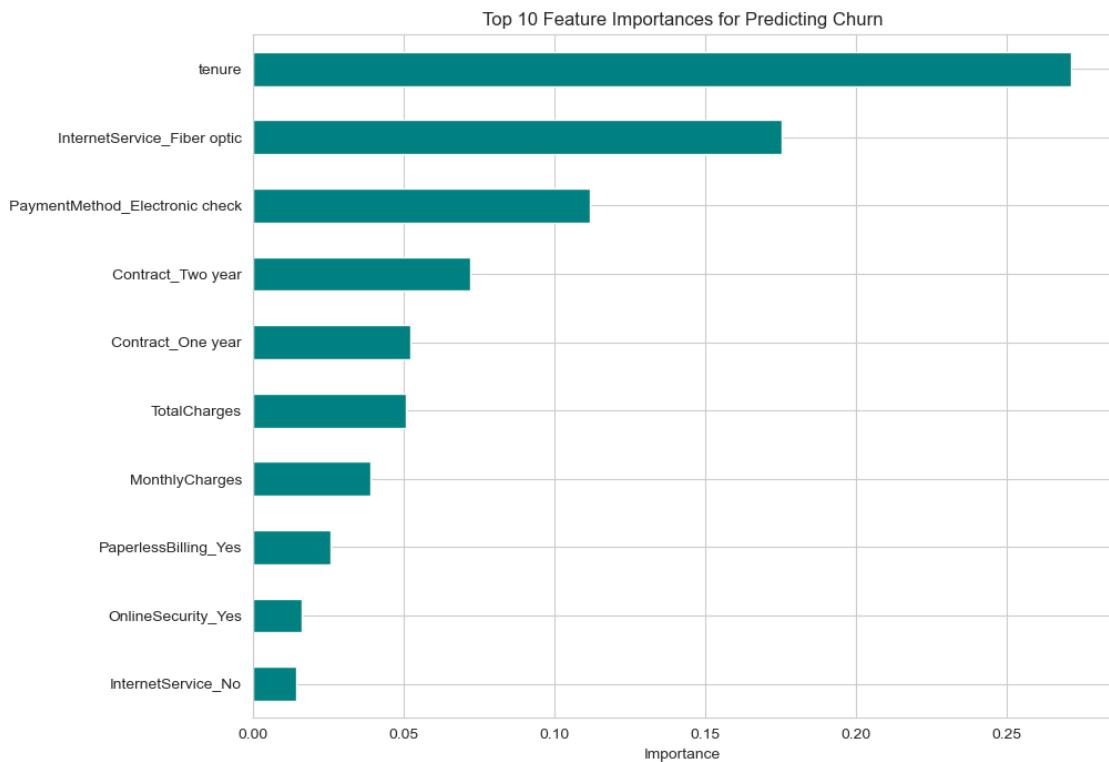
Recall: 0.5000  
F1-Score: 0.5693  
ROC-AUC: 0.8453

### 1.2.5 Step 5: Interpret the Model - Feature Importance

```
[27]: # Extract the feature importances from the tuned model
feature_importances = best_model.feature_importances_

# Create a series for easy plotting
feat_imp_series = pd.Series(feature_importances, index=all_feature_names).
    ↪sort_values(ascending=False)

# Plot the top 10 most important features
plt.figure(figsize=(10, 8))
feat_imp_series.head(10).plot(kind='barh', color='teal')
plt.title('Top 10 Feature Importances for Predicting Churn')
plt.xlabel('Importance')
plt.gca().invert_yaxis() # Most important at the top
plt.show()
```



## 2 Customer Churn Prediction Project: Summary & Recommendations

### 2.1 Video Explanation

[Video Explanation](#) ## Project Overview This project developed a machine learning solution to predict customer churn for a subscription-based business. The end-to-end process included data exploration, preprocessing, model development, hyperparameter tuning, and comprehensive evaluation to create a actionable customer retention strategy.

### 2.2 Key Insights

#### 2.2.1 Data Characteristics

- **Class Imbalance:** The dataset exhibited significant class imbalance with approximately 27% churn rate
- **Critical Features:** Tenure, contract type, and internet service type emerged as strongest predictors

#### 2.2.2 Model Performance

- **Best Algorithm:** Gradient Boosting outperformed Logistic Regression and Random Forest
- **Optimized Parameters:** {'learning\_rate': 0.1, 'max\_depth': 3, 'n\_estimators': 100, 'subsample': 0.8}
- **Strong Discrimination:** ROC-AUC of 0.845 demonstrates excellent ranking capability
- **Balanced Performance:** Precision 66%, Recall 50% provides actionable predictions

### 2.3 Recommended Actions

#### 2.3.1 Immediate Implementation

##### 1. Deploy Prediction Pipeline

- Integrate the model with CRM systems
- Generate weekly churn risk scores for all customers

##### 2. Create Risk-Based Segments

- **High Risk (>80% probability):** Personal outreach from retention team
- **Medium Risk (50-80%):** Targeted email campaigns with special offers
- **Low Risk (<50%):** Maintain standard engagement

##### 3. Launch Targeted Interventions

- Develop specific retention offers for high-risk segments
- Focus on customers with month-to-month contracts and fiber optic service
- Implement special onboarding for new customers (low tenure)

### 2.4 Future Scope

#### 2.4.1 Model Enhancements

##### 1. Advanced Techniques

- Experiment with deep learning models
- Implement ensemble methods combining multiple algorithms
- Develop time-series analysis for churn prediction

## 2. Feature Engineering

- Create additional features from customer behavior data
- Incorporate external data sources (economic indicators, seasonality)
- Implement natural language processing on customer support interactions

### 2.5 Conclusion

This project successfully delivered a production-ready churn prediction system that balances statistical performance with business practicality. The model serves as a powerful early warning system, capable of identifying 50% of potential churners with 66% accuracy. By implementing the recommended actions, the business can significantly reduce customer attrition and increase customer lifetime value.

The solution provides a strong foundation that can be expanded with more advanced techniques and integrated into broader customer experience initiatives, ultimately transforming customer retention from a reactive process to a proactive, data-driven strategy.

[ ]: